




Mark Haylock

<https://github.com/mhaylock>



Faster Pagination with Postgres

(and other SQL servers too!)

Uses of Pagination

Previous / Next

+ Jump to page

Infinite scrolling

Batched iteration through records

Common pagination libraries in Ruby

will_paginate

<https://github.com/ddnexus/pagy>

kaminari

<https://github.com/kaminari/kaminari>

pagy

<https://github.com/ddnexus/pagy>

Common features:

- Paging using OFFSET
- Counting all matching records to calculate a total page count
- Rendering page navigation frontends

Two things that scale terribly in Postgres

Paging using OFFSET
&
COUNT

Apply caution when using these!

OFFSET Pagination

```
SELECT *  
FROM movies  
ORDER BY start_year DESC NULLS LAST, id ASC  
LIMIT 500  
OFFSET 50000; -- Page 101
```


EXPLAIN ANALYZE

```
EXPLAIN (ANALYZE, TIMING OFF) SELECT *  
FROM movies
```

...

Limit (cost=57009.16..57067.49 rows=500 width=122) (actual rows=500 loops=1)

-> Gather Merge (cost=51175.41..106908.17 rows=477676 width=122) (actual rows=50500 loops=1)

Workers Planned: 2

Workers Launched: 2

-> Sort (cost=50175.39..50772.49 rows=238838 width=122) (actual rows=17032 loops=3)

Sort Key: start_year DESC NULLS LAST, id

Sort Method: external merge Disk: 27768kB

Worker 0: Sort Method: external merge Disk: 24472kB

Worker 1: Sort Method: external merge Disk: 24680kB

-> Parallel Seq Scan on movies (cost=0.00..13328.38 rows=238838 width=122) (actual
rows=191071 loops=3)

Planning Time: 0.092 ms

Execution Time: 212.477 ms

Adding an Index

```
CREATE INDEX index_movies_on_start_year_and_id  
ON movies(start_year DESC NULLS LAST, id);
```

Limit (cost=4685.48..4732.33 rows=500 width=122) (actual rows=500 loops=1)

-> Index Scan using index_movies_on_start_year_and_id on movies (cost=0.42..53711.05 rows=573212 width=122) (actual rows=50500 loops=1)

Planning Time: 0.196 ms

Execution Time: 16.226 ms

Seek Method / Keyset Pagination

What if we could help Postgres jump straight to the record after the last record on the previous page?

```
#<Movie:0x00007fb08b205128  
  id: 480934,  
  imdb_title_id: "tt5657280",  
  primary_title: "Viking Destiny",  
  original_title: "Viking Destiny",  
  start_year: 2018,  
  end_year: nil,  
  runtime_minutes: 91,  
  genres: ["Action", "Adventure", "Fantasy"],  
  created_at: Wed, 14 Jul 2021 10:15:46.191499000 UTC +00:00,  
  updated_at: Wed, 14 Jul 2021 10:15:46.191499000 UTC +00:00>
```

```
SELECT *  
FROM movies  
WHERE start_year < 2018  
ORDER BY start_year DESC NULLS LAST, id ASC  
LIMIT 500;
```


Seek Method / Keyset Pagination

What if we could help Postgres jump straight to the record after the last record on the previous page?

```
#<Movie:0x00007fb08b205128
```

```
id: 480934,
```

```
imdb_title_id: "tt5657280",
```

```
primary_title: "Viking Destiny",
```

```
original_title: "Viking Destiny",
```

```
start_year: 2018,
```

```
end_year: nil,
```

```
runtime_minutes: 91,
```

```
genres: ["Action", "Adventure", "Fantasy"],
```

```
created_at: Wed, 14 Jul 2021 10:15:46.191499000 UTC +00:00,
```

```
updated_at: Wed, 14 Jul 2021 10:15:46.191499000 UTC +00:00>
```

```
SELECT *  
FROM movies  
WHERE start_year <= 2018  
      AND NOT (start_year = 2018 AND id <= 480934)  
ORDER BY start_year DESC NULLS LAST, id ASC  
LIMIT 500;
```


EXPLAIN for Seek Method

Limit (cost=0.42..60.14 rows=500 width=122) (actual rows=500 loops=1)

-> Index Scan using index_movies_on_start_year_and_id on movies (cost=0.42..52868.83 rows=442690 width=122) (actual rows=500 loops=1)

Index Cond: (start_year <= 2018)

Filter: ((start_year <> 2018) OR (id < 480934))

Planning Time: 0.122 ms

Execution Time: 0.295 ms

vs the previous EXPLAIN:

Limit (cost=4685.48..4732.33 rows=500 width=122) (actual rows=500 loops=1)

-> Index Scan using index_movies_on_start_year_and_id on movies (cost=0.42..53711.05 rows=573212 width=122) (actual rows=50500 loops=1)

Planning Time: 0.196 ms

Execution Time: 16.226 ms

If column sort order is consistent

Use “Row Values” comparison syntax:

```
SELECT *  
FROM movies  
WHERE (primary_title, id) > ('Viking Destiny', 480934)  
ORDER BY primary_title ASC, id ASC  
LIMIT 500;
```

Limit (cost=0.42..784.47 rows=500 width=122) (actual rows=500 loops=1)

-> Index Scan using index_movies_on_primary_title_and_id on movies (cost=0.42..44927.73 rows=28651 width=122) (actual rows=500 loops=1)

Index Cond: (ROW((primary_title)::text, id) > ROW('Viking Destiny'::text, 480934))

Planning Time: 0.111 ms

Execution Time: 0.567 ms

Going backwards

Means you have to reverse everything,
including the results!

```
SELECT *  
FROM movies  
WHERE (primary_title, id) > ('Viking Destiny', 480934)  
ORDER BY primary_title ASC, id ASC  
LIMIT 500;
```

```
SELECT *  
FROM movies  
WHERE start_year <= 2018  
      AND NOT (start_year = 2018 AND id <= 480934)  
ORDER BY start_year DESC NULLS LAST, id ASC  
LIMIT 500;
```

```
SELECT *  
FROM movies  
WHERE (primary_title, id) < ('Viking Destiny', 480934)  
ORDER BY primary_title DESC, id DESC  
LIMIT 500;
```

```
SELECT *  
FROM movies  
WHERE start_year >= 2018  
      AND NOT (start_year = 2018 AND id >= 480934)  
ORDER BY start_year ASC NULLS FIRST, id DESC  
LIMIT 500;
```


DEMO TIME

Pros & Cons of Seek/Keyset Pagination

Pros:

- Fast
- No page drift as rows are inserted

Cons:

- Every sort order needs an index
- WHERE clause can be complicated
- Going backwards is confusing
- Can't identify current page
- Can't jump to a particular page

Ruby Library

https://github.com/glebm/order_query

References

Pros:

- ***Five ways to paginate in Postgres, from the basic to the exotic***

<https://www.citusdata.com/blog/2016/03/30/five-ways-to-paginate/>

- ***Pagination done the Right way***

<https://www.slideshare.net/MarkusWinand/p2d2-pagination-done-the-postgresql-way>

- ***Paging Through Results***

<https://use-the-index-luke.com/sql/partial-results/fetch-next-page>